

Spatial filtering

Semester 2, 2025

Kris Ehinger

Outline

- Introduction to convolution
- Commonly-used linear filters
- Filters in practice

Learning outcomes

- Explain the convolution and cross-correlation operations
- Identify commonly-used filters and their expected outputs
- Explain practical considerations in implementing filters (efficiency, border handling)

Introduction to convolution

Pixel operator


- Pixel operator: computes an output value at each pixel location, based on the input pixel value

$$g(i, j) = h(f(i, j))$$




Output image g

Input image f

- Example: $g(i, j) = 0.5(f(i, j))$





Gamma correction



Local operator

- Local operator: computes an output value at each pixel location, based on a neighbourhood of pixels around the input pixel
- Example: sharpening filter



Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i, j) = h(u, v) \otimes f(i, j)$$

Output image g Kernel h Input image f

Cross-correlation convolution

$$g(i, j) = \sum_{u, v} f(i + u, j + v) h(u, v)$$


Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i, j) = h(u, v) * f(i, j)$$

Output image g Kernel h Input image f

Convolution operator



$$g(i, j) = \sum_{u, v} f(i - u, j - v) h(u, v)$$

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3		

Output

Position 1:

1x1	5x0
1x0	2x1

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	

Output

Position 2:

5x1	1x0
2x0	8x1

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3

Output

Position 3:

1x1	1x0
8x0	2x1

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1		

Output

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1	3	

Output

Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1	3	16

Output

Cross-correlation vs. convolution

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x,y]$

a	b	c
d	e	f
g	h	i

$H[u,v]$

		i	h	g		
		f	e	d		
		c	b	a		

$F[x,y] \otimes H[u,v]$

Cross-correlation: overlay filter on image

Cross-correlation vs. convolution

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x,y]$

a	b	c
d	e	f
g	h	i

$H[u,v]$

		a	b	c		
		d	e	f		
		g	h	i		

$F[x,y] * H[u,v]$

Convolution: flip filter horizontally and vertically

Summary

- Pixel operator: transform pixel based on its value
- Local operator: transform pixel based on its neighbours
- Convolution (and cross-correlation): operations that apply a linear filter to an image

Common filters

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1


?

(Note that filter sums to 1)

Common filters: Gaussian




$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




Gaussian kernel
Kernel size: 5 x 5 px
 $\sigma = 1$

Blur filters




Gaussian kernel




Average kernel




Common filters: Sobel



Common filters: Sobel





1	2	1
0	0	0
-1	-2	-1



What about colour?

- Consider a 3x4x3 image and 2x2x3 kernel




Convolution output?

Filter size vs. filter support



Convolved with:

$100 \times 100, \sigma = 2$




$100 \times 100, \sigma = 20$



$10 \times 10, \sigma = 2$




?



Designing filters

- How could you detect diagonal edges?



Designing filters

- How could you simulate (linear) motion blur?



Common filters

- Average/blur filters: average pixel values, blur the image
- Sharpening filters: subtract pixel from surround, increase fine detail
- Edge filters: compute difference between pixels, detect oriented edges in image

Filters in practice


Properties of linear filters

- Commutative: $f * h = h * f$
 - Theoretically, no difference between kernel and image
 - But most implementations do care about order
- Associative: $(f * h1) * h2 = f * (h1 * h2)$
 - Usually one option is faster than the other – allows for more efficient implementations
- Distributive over addition
 - $f * (h1 + h2) = (f * h1) + (f * h2)$
- Multiplication cancels out
 - $kf * h = f * kh = k(f * h)$

Efficient filtering


- Multiple filters: generally more efficient to combine 2D filters ($h1 * h2 * h3...$) and filter image just once





*  =
Gaussian
blur filter



*  =
Horizontal
derivative
filter



*  =
Derivative-of-
Gaussian filter



Efficient filtering

- Separable filters: generally more efficient to filter with two 1D filters than one 2D filter
- For example, the 2D Gaussian can be expressed as a product of two 1D Gaussians (in x and y)

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

Separable filters

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors into a product
of 1D filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution along
rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix} = \begin{bmatrix} & & \\ & 65 & \\ & & \end{bmatrix}$$

Convolution output size

- Valid convolution: the output image is smaller than the input image
- Why?

0	0	0	0	0	0	0
0	0	255	255	255	0	0
0	255	0	0	0	255	0
0	255	0	0	0	0	0
0	255	0	0	0	255	0
0	0	255	255	255	0	0
0	0	0	0	0	0	0

 $*$


$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

 $=$

57	85	85	85	57
85	113	85	85	57
85	85	0	57	57
85	113	85	85	57
57	85	85	85	57

Border handling

- Pad with constant value



Filter: Gaussian blur

Border handling

- Wrap image



Filter: Gaussian blur



Border handling

- Clamp / replicate the border value



Filter: Gaussian blur




Border handling

- Reflect image



Filter: Gaussian blur



Practical considerations

- Think about how to implement filters efficiently
 - Images are big, so efficient filtering can save a lot of time!
- Think about how to handle borders
 - No one-size-fits-all solution
 - Wrap is ideal for tiling textures (but not photos)
 - Clamp/replicate tends to work well for photos

Summary

- Linear filters: first step of almost all computer vision systems
- Linear filters are just a first step – you can't build complex feature detectors from just linear filters